

Modelling and Analysis of Data

Exam no. 39

22. januar 2022

Indhold

1 Problem 1 - Maximum Likelihood Estimation	1
1.1	1
2 Problem 2 - Statistics	3
3 Problem 3 - Principal Component Analysis	3
3.1 1) - Mean point of the data	3
3.2 2) - Two eigenvalues	4
3.3 3) - Two eigenvectors	4
4 Problem 4 - Regression	5
4.1 a)	5
4.2 b)	5
4.3 c)	6
5 Problem 5 - Random Forests	7
5.1 Finding Thresholds	7
6 Problem 6 - Classification & Validation	10
6.1 a) - Implement random forest training	10
6.2 b) - Finding the optimal set of random forest classifier parameters . . .	10
6.3 c) - My results	11
6.4 Code for a) & b)	11
7 Problem 7 - K-means Clustering & Principal Component Analysis	14
7.1 a)	14
7.2 b)	14
7.3 c)	15
7.4 d)	15
7.5 e)	16
8 Appendix A	17
8.1	17
9 Appendix B	18
9.1	18
10 Appendix Q6	19
11 Appendix X	22
11.1 Q6_c	22

1 Problem 1 - Maximum Likelihood Estimation

1.1

Assessing the following:

$$f(x) \begin{cases} \frac{1}{\sigma\sqrt{2\pi}} \cdot \frac{1}{x} \cdot \exp\left(-\frac{1}{2}\left(\frac{\log(x)-\mu}{\sigma}\right)^2\right) & \text{for } x \in \mathbb{R}_+ \\ 0 & \text{otherwise} \end{cases}$$

Since I can fix the parameter σ I can express the likelihood L of observing the given dataset x_1, \dots, x_N of N i.i.d samples from X as the joint distribution all the N i.i.d samples conditioning on μ as:

$$L = p(x_1, \dots, x_N | \mu) = \prod_{n=1}^N p(x_n | \mu) = \prod_{n=1}^N f(x_n; \mu)$$

The PDF is now considered a function of both x_n and μ I'm going to make take the logarithm of the expression, to make it more manageable, without changing the estimate. When trying to find the optimal $\hat{\mu}$ that maximizes the likelihood L .

$$\begin{aligned} \log L &= \log\left(\prod_{n=1}^N f(x_n; \mu)\right) \\ &= \sum_{n=1}^N \log(f(x_n; \mu)) \\ &= \sum_{n=1}^N \left[\log\left(\frac{1}{\sigma\sqrt{2\pi}}\right) + \log\left(\frac{1}{x_n}\right) + \log\left(\exp\left(-\frac{1}{2}\left(\frac{\log(x)-\mu}{\sigma}\right)^2\right)\right) \right] \\ &= \sum_{n=1}^N \left[\log\left(\frac{1}{\sigma\sqrt{2\pi}}\right) + \log\left(\frac{1}{x_n}\right) - \frac{1}{2}\left(\frac{\log(x)-\mu}{\sigma}\right)^2 \right] \end{aligned}$$

To find the MLE $\hat{\mu}$ for the parameter μ I optimize L by solving $\frac{\partial \log L}{\partial \mu} = 0$ When taking derivatives I'm only interested in terms that involve μ and so I have the following expression.

$$\begin{aligned} \frac{\partial \log L}{\partial \mu} &= \frac{\partial}{\partial \mu} \left[\sum_{n=1}^N -\frac{1}{2} \left(\frac{\log(x)-\mu}{\sigma} \right)^2 \right] \\ &= \sum_{n=1}^N -\left(\frac{\log(x)-\mu}{\sigma} \right) \cdot \left[\frac{\log(x)-\mu}{\sigma} \cdot \frac{\partial}{\partial \mu} \right] \quad (\text{Chain rule}) \\ &= \sum_{n=1}^N -\left(\frac{\log(x)-\mu}{\sigma} \right) \cdot \left(-\frac{1}{\sigma} \right) \\ &= \frac{1}{\sigma^2} \sum_{n=1}^N \log(x_n) - \mu \end{aligned}$$

solving $\frac{\partial \log L}{\partial \mu} = 0$ with respect to μ

$$\begin{aligned}
 0 &= \frac{1}{\sigma^2} \sum_{n=1}^N \log(x_n) - \mu = \sum_{n=1}^N \log(x_n) - \mu = \left[\sum_{n=1}^N \log(x_n) \right] - N\mu \\
 &\Downarrow \\
 N\mu &= \sum_{n=1}^N \log(x_n) \\
 &\Downarrow \\
 \mu &= \frac{1}{N} \sum_{n=1}^N \log(x_n) = \frac{1}{N} \log(x_1 \cdot x_2 \cdot \dots \cdot x_N) = \log\left((x_1 \cdot x_2 \cdot \dots \cdot x_N)^{\frac{1}{N}}\right) \\
 &= \log(\sqrt[N]{x_1 \cdot x_2 \cdot \dots \cdot x_N})
 \end{aligned}$$

The above calculations is a necessary condition in order to find an optimal solution. Furthermore, to check if I'm dealing with a global maximum or a global minimum I have to take the second derivative of my function.

$$\frac{\partial^2 \log L}{\partial \mu^2} = -\frac{N}{\sigma^2}$$

Assessing the result from the second derivative, I can see that my function will always be negative. Thus by the second derivative test

$$\hat{\mu} = \log(\sqrt[N]{x_1 \cdot x_2 \cdot \dots \cdot x_N})$$

is a global maximum

2 Problem 2 - Statistics

3 Problem 3 - Principal Component Analysis

3.1 1) - Mean point of the data

Step one is to calculate the mean point of the data with respect to both x and y

$$\bar{x} = \frac{1}{8}(0.6 + 0.5 + 1.1 + (-0.5) + 0.8 + 0.2 + (-0.1) + 1.0) = 0.4500$$

$$\bar{y} = \frac{1}{8}(1.1 + 1.0 + 2.0 + 0.2 + (-0.1) + (-0.1) + (-1.5) + 2.5) = 0.6375$$

Step two is to normalize the dataset. This is done by taking each original point and subtract the mean. $p - [\bar{x}, \bar{y}]$

Normalizing each point

$$p1 = (0.6, 1.1) - [0.4500, 0.6375] = [0.1500, 0.4625] \quad (1)$$

$$p2 = (0.5, 1.0) - [0.4500, 0.6375] = [0.0500, 0.3625] \quad (2)$$

$$p3 = (1.1, 2.0) - [0.4500, 0.6375] = [0.6500, 1.3625] \quad (3)$$

$$p4 = (-0.5, 0.2) - [0.4500, 0.6375] = [-0.9500, -0.4375] \quad (4)$$

$$p5 = (0.8, -0.1) - [0.4500, 0.6375] = [0.3500, -0.7375] \quad (5)$$

$$p6 = (0.2, -0.1) - [0.4500, 0.6375] = [-0.2500, -0.7375] \quad (6)$$

$$p7 = (-0.1, -1.5) - [0.4500, 0.6375] = [-0.5500, -2.1375] \quad (7)$$

$$p8 = (1.0, 2.5) - [0.4500, 0.6375] = [0.5500, 1.8625] \quad (8)$$

Step three is to calculate the covariances for XX, XY and YY using the normalized values for p1-p8

$$\begin{aligned} cov(xx) = \frac{1}{8} \Big(& 0.15^2 + 0.05^2 + 0.65^2 + (-0.95)^2 \\ & + (-0.25)^2 + (-0.55)^2 + 0.55^2 \Big) = 0.2675 \end{aligned} \quad (9)$$

$$\begin{aligned} cov(yy) = \frac{1}{8} \Big(& 0.4625^2 + 0.3625^2 + 1.3625^2 + (-0.4375)^2 + \\ & (-0.7375)^2 + (-0.7375)^2 + (-2.1375)^2 + 1.8625^2 \Big) = 1.439843750 \end{aligned} \quad (10)$$

$$\begin{aligned} cov(xy) = \frac{1}{8} \Big(& (0.15 \cdot 0.4625 + 0.05 \cdot 0.3625 + 0.65 \cdot 1.3625 + ((-0.95) \cdot (-0.4375)) \\ & + 0.35 \cdot (-0.7375) + ((-0.25) \cdot (-0.7375)) + ((-0.55) \cdot (-2.1375)) \\ & + 0.55 \cdot 1.8625) \Big) = 0.439375 \end{aligned} \quad (11)$$

having done the three calculations I can now form the covariance matrix inserting the calculated values into their respectable place:

$$\begin{aligned} cov &= \begin{pmatrix} xx & xy \\ xy & yy \end{pmatrix} = \\ cov &= \begin{pmatrix} 0.2675 & 0.439375 \\ 0.439375 & 1.439843750 \end{pmatrix} \end{aligned}$$

3.2 2) - Two eigenvalues

In order to find the eigenvalues and the eigenvectors I need to solve **Characteristic equation** of the covariance matrix. The Characteristic equation is denoted as:

$$\det(A - \lambda I) = 0$$

To calculate the determinant you calculate the following:

$$\text{determinant} = \begin{vmatrix} a & b \\ c & c \end{vmatrix} = ad - bc$$

Solving the system equation that I get from using the Characteristics equation on the covariance matrix I get the following eigenvalues.

$$\begin{vmatrix} 0.2675 - \lambda & 0.439375 \\ 0.439375 & 1.439843750 - \lambda \end{vmatrix} = 0$$

First I get the equation by calculating the determinant

$$\det(cov) \rightarrow ((0.2675 - \lambda) \cdot (1.439843750 - \lambda)) - (0.439375 \cdot 0.439375)$$

Then setting the equation equal to zero and solving it will give me the two values for λ which is the eigenvalues.

$$((0.2675 - \lambda) \cdot (1.439843750 - \lambda)) - (0.439375 \cdot 0.439375) = 0$$

$$\text{eigenvalue}_1 = 0.1211093468 \quad \text{eigenvalue}_2 = 1.586234403$$

3.3 3) - Two eigenvectors

In order to find the eigenvectors I use the definition of an eigenvector $A\bar{x} = \lambda\bar{x}$ and solve this equation.

$$cov \cdot \bar{x} = \lambda \cdot \bar{x}$$

I'm going to do the calculations in Maple since I'm dealing with odd numbers. I also going to reduce some of the decimals by rounding off just for readability sake.

$$\text{eigenvector}_1 = [-0.9487, -0.3161] \quad \text{eigenvector}_2 = [0.3161, -0.9487]$$

note that signs of the eigenvector can be different. e.g. eigenvector_1 could be $[0.9487, 0.3161]$

4 Problem 4 - Regression

4.1 a)

Using the euclidean distance I get an RSME value of

$$RMSE = 0.82430$$





Figure 1: Plot of kNN with $k = 3$

4.2 b)

Using the other distance calculation as given in the exam set I get the following RSME value

$$RMSE = 1.100519$$



Figur 2: Plot of kNN with $k = 3$ and distance matrix

The distance matrix will lower/lessen the influence of all the features/attributes except for the last two, since they are set to 1.0.

4.3 c)

5 Problem 5 - Random Forests

5.1 Finding Thresholds

There are four places that I should take into consideration as possible thresholds.

option	x_i	x_{i+1}
1	20	30
2	60	70
3	90	100
4	100	110

There are four positions that could be possible threshold, however, I'm going to ignore option 4, since it will be the same as option 1, except with reversed values. These four places have been found using the same approach as shown in our lecture, and lecture video (L10_Classification_Regression_2_Part3, starting at 17min 50sec.). This means, that I'm only going to look at positions where there is a "change" in labels for a given datapoint compared to its predecessor or successor. Since there is no reason to look at possible thresholds where the neighbors to a given feature have the same labels.

This process can be formalized:

$$x_i + \left(\frac{x_{i+1} - x_i}{2}\right) : y_{i+1} \neq y_i$$

As stated, option 4 is the same as option 1, so I will only perform the calculations for option 1-3.

$$\text{Option 1} = 20 + \frac{(30 - 20)}{2} = 25$$

$$\text{Option 2} = 60 + \frac{(70 - 60)}{2} = 65$$

$$\text{Option 3} = 90 + \frac{(100 - 90)}{2} = 95$$

Using my calculated values for possible thresholds I have three different decision trees to take into consideration.

Tree 1:






Figure 3: Decision tree with threshold = 25

Tree 2:



Figur 4: Decision tree with threshold = 65

Tree 3:



Figur 5: Decision tree with threshold = 95

From here, I will start by calculating the Entropy values for splits in my trees.
Aswell as the information gain using the functions given in our lecture slides (L10_Classification_Regression_2_Part3 p. 8)

$$\text{Entropy } H(T) = -p_+ \cdot \log_2 p_+ - p_- \cdot \log_2 p_-$$

$$\text{Gain}(T, j) = H(T) - \sum_{i=0,1} \frac{|T_i|}{|T|} H(T_i)$$

Calculating the information gain for tree 1 where the threshold is = 25:

$$\begin{aligned} \text{Gain}(T, j) &= H(5 \text{ pos}, 7 \text{ neg}) - \frac{10}{12} H(5 \text{ pos}, 5 \text{ neg}) - \frac{2}{12} H(0 \text{ pos}, 2 \text{ neg}) \\ &= 0.979 - \frac{10}{12} 1 - \frac{2}{12} 0 \approx 0.145 \end{aligned}$$

Calculating the information gain for tree 2 where the threshold is = 65:

$$\begin{aligned} \text{Gain}(T, j) &= H(5 \text{ pos}, 7 \text{ neg}) - \frac{6}{12} H(4 \text{ pos}, 2 \text{ neg}) - \frac{6}{12} H(1 \text{ pos}, 5 \text{ neg}) \\ &= 0.979 - \frac{10}{12} 0.918 - \frac{2}{12} 0.650 \approx 0.195 \end{aligned}$$

Calculating the information gain for tree 3 where the threshold is = 95:

$$\begin{aligned} \text{Gain}(T, j) &= H(5 \text{ pos}, 7 \text{ neg}) - \frac{6}{12} H(4 \text{ pos}, 2 \text{ neg}) - \frac{6}{12} H(1 \text{ pos}, 5 \text{ neg}) \\ &= 0.979 - \frac{10}{12} 0.918 - \frac{2}{12} 0.991 \approx 0.0062 \end{aligned}$$

Since the goal is to find the tree which gives us the highest information gain. I can see that the second decision tree with a threshold of 65 is the best one.

Thus, the optimal threshold is 65.

6 Problem 6 - Classification & Validation

NOTE that I've included the full python code in the appendix

6.1 a) - Implement random forest training

I've implemented the Random Forest Training using the existing library from *Skikit-learn*, more specific the package: `sklearn.ensemble.RandomForestClassifier`. Below I've added the most notable part from my code in order to load the data. I'm reusing the "loaddata" function that's been handed out, though I've modified it slightly in line 12 in the code snippet.

```
1  # Modified the loaddata function to make the training features
   ↳ into 2d np.array, more specifically the 't value'
2  def loaddata(filename):
3      """Load the balloon data set from filename and return t, X
4          t - N-dim. vector of target (temperature) values
5          X - N-dim. vector containing the inputs (lift) x for
   ↳ each data point
6      """
7      # Load data set from CSV file
8      Xt = np.loadtxt(filename, delimiter=',')
9
10     # Split into data matrix and target vector
11     X = Xt[:,0] # Matrix
12     t = Xt[:,1:] # Vector
13
14     return t, X
15
16 # Loads the training data
17 X_train, t_train =
   ↳ loaddata("./accent-mfcc-data_shuffled_train.txt")
18
19 # Loads the validation data
20 X_validation, t_validation =
   ↳ loaddata("./accent-mfcc-data_shuffled_validation.txt")
```

To see how I've chosen to set up the Random Forest Classifier, training the algorithm using the it to predict - see the submitted code in section 6.4

6.2 b) - Finding the optimal set of random forest classifier parameters

I've chosen to use two extra packages in order to implement my Random Forest algorithm. It hasn't been specified that I'm not allowed to use other packages from the **Sklean** library so I assume that this is a possible solution.

I've made use of **ParameterGrid** from the *sklearn.model_selection* and **accuracy_score** from the *sklearn.metrics*.

I'm using ParameterGrid to generate a list of all the possible permutations of the specified parameters used for a given iteration to perform the optimized search. With each iteration the parameters and the resulting performance metrics are added to a list which can be sorted easily. I'm currently using the accuracy_score as the first metric.

see the submitted code in section 6.4

6.3 c) - My results

Looking at the results, it's possible to see that the performance of the algorithm increases when tree depth, number of features that's taken into consideration and the complexity of the criterion is also increased. However it comes with the cost of a more computational complexity aswell. The second listing in section 6.4 shows the terminal output from the optimization loop. I've appended a sorted print of the permuted parameters in the appendix.

Assessing the result, I get the optimal results from the algorithm with the parameters:

```
{'criterion': 'entropy', 'max_tree_depth': 10, 'max_features': 'sqrt'}
```

6.4 Code for a) & b)

testtesttest

```
1  # Adding the parameters to test given in the exam question.
   ↳ Using these to find the optimal set of random forest
   ↳ classifier parameters.
2  random_f_parameters = {
3      'criterion'      : ['gini', 'entropy'],
4      'max_tree_depth' : [2,5,6,10,15],
5      'max_features'   : ['sqrt', 'log2']
6  }
7
8  # Empty array for the result metrics
9  res = np.empty((0,3))
10
11 # Looping through the chosen(given) parameters and setting up
   ↳ the Random forest classifier each time.
12 for params in list(ParameterGrid(random_f_parameters)):
13     clf = RandomForestClassifier(
14         criterion      = params['criterion'],
15         max_depth      = params['max_tree_depth'],
16         max_features   = params['max_features'])
17
```

```
18     # Training using the created classifier with the given
19     ↪ parameters.
20     clf.fit(X_train, t_train)
21
22     # number of correctly classified validation samples
23     t_prediction = clf.predict(X_validation)
24     acc_score = accuracy_score(t_validation, t_prediction)
25
26     # probability associated with classification
27     t_probability = clf.predict_proba(X_validation)
28     probability_score = np.mean([t_probability[int(t_val)]
29     ↪ for (t_probability, t_val)
30     ↪ in zip(t_probability,
31     ↪ ↪ t_validation)])
32
33     print("Accuracy score: %.2f"
34     ↪ %acc_score)
35     print("Average probability assigned to correct classes:
36     ↪ %.2f"
37     ↪ %probability_score)
38
39     # print the parameters if new ones are more optimal than
40     ↪ previously tried ones.
41     if len(res) > 0 and (acc_score > res[-1,1]
42     ↪ or (acc_score == res[-1,1]
43     ↪ and probability_score >
44     ↪ ↪ res[-1,2])):
45         print(params)
46
47     # accumulate results
48     res = np.append(res, np.array([[params, acc_score,
49     ↪ probability_score]]), axis=0)
50     # sort the results in ascending order
51     res = res[np.lexsort((res[:,1], res[:,2]))]
```

Accuracy score: 0.48
Average probability assigned to correct classes: 0.36
Accuracy score: 0.66
Average probability assigned to correct classes: 0.48
{'criterion': 'gini', 'max_features': 'sqrt', 'max_tree_depth': 5}
Accuracy score: 0.73
Average probability assigned to correct classes: 0.51
{'criterion': 'gini', 'max_features': 'sqrt', 'max_tree_depth': 6}
Accuracy score: 0.77
Average probability assigned to correct classes: 0.56
{'criterion': 'gini', 'max_features': 'sqrt', 'max_tree_depth': 10}
Accuracy score: 0.77
Average probability assigned to correct classes: 0.56
{'criterion': 'gini', 'max_features': 'sqrt', 'max_tree_depth': 15}
Accuracy score: 0.48
Average probability assigned to correct classes: 0.36
Accuracy score: 0.69
Average probability assigned to correct classes: 0.48
Accuracy score: 0.71
Average probability assigned to correct classes: 0.51
Accuracy score: 0.82
Average probability assigned to correct classes: 0.56
{'criterion': 'gini', 'max_features': 'log2', 'max_tree_depth': 10}
Accuracy score: 0.79
Average probability assigned to correct classes: 0.57
{'criterion': 'gini', 'max_features': 'log2', 'max_tree_depth': 15}
Accuracy score: 0.52
Average probability assigned to correct classes: 0.37
Accuracy score: 0.73
Average probability assigned to correct classes: 0.51
Accuracy score: 0.74
Average probability assigned to correct classes: 0.54
Accuracy score: 0.81
Average probability assigned to correct classes: 0.57
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 10}
Accuracy score: 0.81
Average probability assigned to correct classes: 0.58
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 15}
Accuracy score: 0.51
Average probability assigned to correct classes: 0.38
Accuracy score: 0.75
Average probability assigned to correct classes: 0.51
Accuracy score: 0.78
Average probability assigned to correct classes: 0.54
Accuracy score: 0.81
Average probability assigned to correct classes: 0.58
Accuracy score: 0.78
Average probability assigned to correct classes: 0.57

Side 13 of 16

Figur 6: Caption of figure here

7 Problem 7 - K-means Clustering & Principal Component Analysis

7.1 a)

I've decided to make my own function to normalize my data. It normalizes the input value x by subtracting the mean and division by the standard deviation.

```
1  def normalizeData(x):  
2      '''  
3      Normalizes the input using mean and the standard deviation.  
4      Normalization is done on all the columns in the 2d  
5      ↪ numpy-array 'x'  
6  
7      Returns the normalized 'x'  
8      '''  
9      return (x - np.mean(x, axis=0)) / np.std(x, axis=0)
```

7.2 b)

Below is a codesnippet of my implementation of the k-mean clustering. I've defined multiple helperfunctions, which can be found in the source code or in the appendix.

```
1  def k_mean_clustering(k, data, centroids):  
2      """  
3      K-means clustering  
4  
5      Params:  
6      -----  
7      k : number of clusters  
8      data : (n_samples, n_features) (normalized) data matrix  
9  
10     Returns:  
11     -----  
12     assignments, intra_cluster_dist : tuple  
13  
14     """  
15     current_assignments = assign_datapoints_to_centroids(data,  
16     ↪ centroids)  
17  
18     new_assignments = [] # initial empty value  
19  
20     while not np.array_equal(current_assignments,  
21     ↪ new_assignments):
```



```
20     # repeat until assignments does not change
21     centroids = calculate_new_centroids(data,
22         ↪ current_assignments, centroids)
23     current_assignments = new_assignments
24     new_assignments = assign_datapoints_to_centroids(data,
25         ↪ centroids)
26
27     intra_cluster_dist = compute_sum_intra_cluster_dist(data,
28         ↪ current_assignments, centroids)
29
30     return current_assignments, intra_cluster_dist
```

The results of the k-means cluster algorithm with $k = 3$. The algorithm has been implemented with `np.random.default_rng` which constructs a random generator so that a desired "seed_value" can be reproduced.

7.3 c)

Here's a print of the values that I get from the k-means clustering algorithm, when using seed: 1

```
Calculations with seed value: 1
Smallest Intra-Cluster Distance: 268.5510
Number of samples in cluster 0: 67
Number of samples in cluster 1: 66
Number of samples in cluster 2: 67
```

7.4 d)

I've reused code that has been handed out in the assignments aswell as my own implementation of PCA from Assignment 5.

7.5 e)




Figure 7: Caption of figure here

8 Appendix A

8.1 ...

```
1 def minted_for_code_here
```

9 Appendix B

9.1 ...

```
1 def minted_for_code_here
```

10 Appendix Q6

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.metrics import accuracy_score
4 from sklearn.model_selection import ParameterGrid
5 from sklearn.ensemble import RandomForestClassifier
6
7 # Modified the loaddata function to make the training features
8   ↳ into
9 # 2d np.array, more specifically the 't value'
10 def loaddata(filename):
11     """Load the balloon data set from filename and return t, X
12         t - N-dim. vector of target (temperature) values
13         X - N-dim. vector containing the inputs (lift) x for
14         ↳ each data point
15     """
16     # Load data set from CSV file
17     Xt = np.loadtxt(filename, delimiter=',')
18
19     # Split into data matrix and target vector
20     X = Xt[:,0] # Matrix
21     t = Xt[:,1:] # Vector
22
23     return t, X
24
25 '''
26 QUESTION 6 - A
27 '''
28 # Loads the training data
29 X_train, t_train =
30   ↳ loaddata("./accent-mfcc-data_shuffled_train.txt")
31
32 # Loads the validation data
33 X_validation, t_validation =
34   ↳ loaddata("./accent-mfcc-data_shuffled_validation.txt")
35
36 print("Shape of train targets: %s" %str(t_train.shape))
37 print("Shape of train features: %s" %str(X_train.shape))
38 print("Shape of val targets: %s" %str(t_validation.shape))
39 print("Shape of val features: %s" %str(X_validation.shape))
40 print()
41
42 '''
```

```
40 QUESTION 6 - B
41 '''
42 # Adding the parameters to test given in the exam question.
43 # Using these to find the optimal set of random forest
44 ↪ classifier parameters.
45 random_f_parameters = {
46     'criterion'      : ['gini', 'entropy'],
47     'max_tree_depth' : [2,5,6,10,15],
48     'max_features'   : ['sqrt', 'log2']
49 }
50
51 # Empty array for the result metrics
52 res = np.empty((0,3))
53
54 # Looping through the chosen(given) parameters and setting up
55 ↪ the
56 # Random forest classifier each time.
57 for params in list(ParameterGrid(random_f_parameters)):
58     clf = RandomForestClassifier(
59         criterion      = params['criterion'],
60         max_depth      = params['max_tree_depth'],
61         max_features    = params['max_features'])
62
63     # Training using the created classifier with the given
64     ↪ parameters.
65     clf.fit(X_train, t_train)
66
67     # number of correctly classified validation samples
68     t_prediction = clf.predict(X_validation)
69     acc_score = accuracy_score(t_validation, t_prediction)
70
71     # probability associated with classification
72     t_probability = clf.predict_proba(X_validation)
73     probability_score = np.mean([t_probability[int(t_val)]
74                                 for (t_probability, t_val)
75                                 in zip(t_probability, t_validation)])
76
77     print("Accuracy score: %.2f"
78           %acc_score)
79     print("Average probability assigned to correct classes:
80           ↪ %.2f"
81           %probability_score)
82
83     # print the parameters if new ones are more optimal
84     # than previously tried ones.
```

```
81     if len(res) > 0 and (acc_score > res[-1,1]
82                           or (acc_score == res[-1,1]
83                               and probability_score >
84                                 res[-1,2])):
85
86         print(params)
87
88     # accumulate results
89     res = np.append(res, np.array([[params, acc_score,
90                                     probability_score]]), axis=0)
91     # sort the results in ascending order
92     res = res[np.lexsort((res[:,1], res[:,2]))]
```

```
91 for x in res:
92     print(x[0])
```

11 Appendix X

11.1 Q6_c

This is a sorted print of the permuted parameters from the Random Forest algorithm, which yields matrices that are increasingly better.

I've rearranged the terminal output slightly for better readability.

```
{'criterion': 'gini'      , 'max_features': 'sqrt', 'max_tree_depth': 2}
{'criterion': 'gini'      , 'max_features': 'log2', 'max_tree_depth': 2}
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 2}
{'criterion': 'entropy', 'max_features': 'log2', 'max_tree_depth': 2}
{'criterion': 'gini'      , 'max_features': 'sqrt', 'max_tree_depth': 5}
{'criterion': 'gini'      , 'max_features': 'log2', 'max_tree_depth': 5}
{'criterion': 'entropy', 'max_features': 'log2', 'max_tree_depth': 5}
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 5}
{'criterion': 'gini'      , 'max_features': 'log2', 'max_tree_depth': 6}
{'criterion': 'gini'      , 'max_features': 'sqrt', 'max_tree_depth': 6}
{'criterion': 'entropy', 'max_features': 'log2', 'max_tree_depth': 6}
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 6}
{'criterion': 'gini'      , 'max_features': 'sqrt', 'max_tree_depth': 10}
{'criterion': 'gini'      , 'max_features': 'log2', 'max_tree_depth': 10}
{'criterion': 'gini'      , 'max_features': 'sqrt', 'max_tree_depth': 15}
{'criterion': 'gini'      , 'max_features': 'log2', 'max_tree_depth': 15}
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 10}
{'criterion': 'entropy', 'max_features': 'log2', 'max_tree_depth': 15}
{'criterion': 'entropy', 'max_features': 'log2', 'max_tree_depth': 10}
{'criterion': 'entropy', 'max_features': 'sqrt', 'max_tree_depth': 15}
```