



Faculty of Science

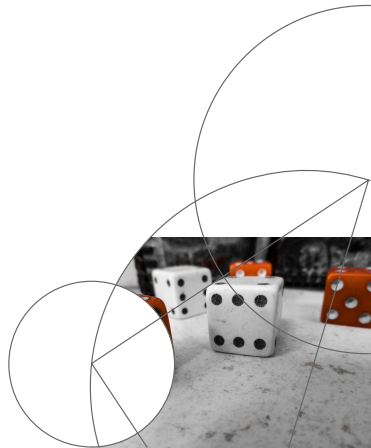
L3 – Non-Linear Regression

Modelling and Analysis of Data

Fabian Gieseke

Machine Learning Section
Department of Computer Science
University of Copenhagen

Universitetsparken 1, 1-1-N110
fabian.gieseke@di.ku.dk



Outline

- 1 Recap: Linear Regression
- 2 Non-Linear Response, Overfitting, and Cross-Validation
- 3 Regularisation
- 4 Summary & Outlook

Outline

- 1 Recap: Linear Regression
- 2 Non-Linear Response, Overfitting, and Cross-Validation
- 3 Regularisation
- 4 Summary & Outlook

Multivariate Linear Regression

- Given: Pairs of the form $(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N) \in \mathbb{R}^D \times \mathbb{R}$.
- Let's "augment" all data points $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$. This yields an **augmented data matrix** $\mathbf{X} \in \mathbb{R}^{N \times (D+1)}$ and an associated target vector $\mathbf{t} \in \mathbb{R}^N$:

$$\mathbf{X} = \begin{bmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,D} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,D} \\ \vdots & & & & \\ 1 & x_{N,1} & x_{N,2} & \dots & x_{N,D} \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- As before, we can write the overall loss in the following form:

Overall Loss

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N (f(\mathbf{x}_n; \mathbf{w}) - t_n)^2 = \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{t})^T (\mathbf{X}\mathbf{w} - \mathbf{t})$$

Summary: Multivariate Linear Regression

- Given: Pairs of the form $(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N) \in \mathbb{R}^D \times \mathbb{R}$.
- Goal: Find $(D+1)$ -dimensional weight vector $\hat{\mathbf{w}} = [\hat{w}_0, \hat{w}_1, \dots, \hat{w}_D]^T$ that minimizes $\mathcal{L}(\mathbf{w}) = \frac{1}{N}(\mathbf{X}\mathbf{w} - \mathbf{t})^T(\mathbf{X}\mathbf{w} - \mathbf{t})$, i.e., which is a solution for

$$\begin{aligned} \nabla \mathcal{L}(\mathbf{w}) &= \mathbf{0} \\ \Leftrightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned} \tag{1}$$

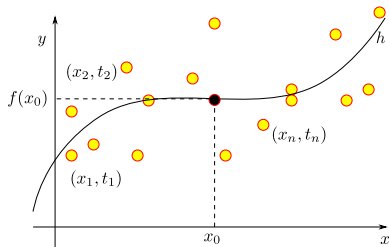
Computation in Practice

- Definition of data matrix $\mathbf{X} \in \mathbb{R}^{N \times (D+1)}$
(make use of Numpy arrays and functions!)
- There are different ways to compute an optimal weight vector $\hat{\mathbf{w}}$:
 - Compute $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$ (e.g., via `numpy.linalg.inv`)
 - Directly solve system of equations (1) (e.g., via `numpy.linalg.solve`)
 - ...
- For new point $\mathbf{x}_{new} \in \mathbb{R}^D$: Compute $t_{new} = [1, \mathbf{x}_{new}^T] \hat{\mathbf{w}}$

Outline

- 1 Recap: Linear Regression
- 2 Non-Linear Response, Overfitting, and Cross-Validation
- 3 Regularisation
- 4 Summary & Outlook

“Non-Linear” Models?



Quadratic Models

- Let's focus again on $D = 1$, i.e., on input data of the form $x_n \in \mathbb{R}$.

Quadratic Models

- Let's focus again on $D = 1$, i.e., on input data of the form $x_n \in \mathbb{R}$.
- Now, let's "augment" all data points x_1, x_2, \dots, x_N , now with an additional column containing x_n^2 . This yields an **augmented data matrix** $\mathbf{X} \in \mathbb{R}^{N \times 3}$ and an associated target vector $\mathbf{t} \in \mathbb{R}^N$:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & & \\ 1 & x_N & x_N^2 \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

Quadratic Models

- Let's focus again on $D = 1$, i.e., on input data of the form $x_n \in \mathbb{R}$.
- Now, let's "augment" all data points x_1, x_2, \dots, x_N , now with an additional column containing x_n^2 . This yields an **augmented data matrix** $\mathbf{X} \in \mathbb{R}^{N \times 3}$ and an associated target vector $\mathbf{t} \in \mathbb{R}^N$:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & & \\ 1 & x_N & x_N^2 \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- Note: Basically as before, just a "new" input feature/variable.
(i.e., we have generated a new column based on an existing column)

Quadratic Models

- Let's focus again on $D = 1$, i.e., on input data of the form $x_n \in \mathbb{R}$.
- Now, let's "augment" all data points x_1, x_2, \dots, x_N , now with an additional column containing x_n^2 . This yields an **augmented data matrix** $\mathbf{X} \in \mathbb{R}^{N \times 3}$ and an associated target vector $\mathbf{t} \in \mathbb{R}^N$:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & & \\ 1 & x_N & x_N^2 \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- Note: Basically as before, just a "new" input feature/variable.
(i.e., we have generated a new column based on an existing column)
- As before: $f(\mathbf{x}; \mathbf{w}) = \mathbf{x}^T \mathbf{w}$ with $\mathbf{x} = [1, x, x^2]^T$ and $\mathbf{w} = [w_0, w_1, w_2]^T$

Quadratic Models

- Let's focus again on $D = 1$, i.e., on input data of the form $x_n \in \mathbb{R}$.
- Now, let's “augment” all data points x_1, x_2, \dots, x_N , now with an additional column containing x_n^2 . This yields an **augmented data matrix** $\mathbf{X} \in \mathbb{R}^{N \times 3}$ and an associated target vector $\mathbf{t} \in \mathbb{R}^N$:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & & \\ 1 & x_N & x_N^2 \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- Note: Basically as before, just a “new” input feature/variable.
(i.e., we have generated a new column based on an existing column)
- As before: $f(\mathbf{x}; \mathbf{w}) = \mathbf{x}^T \mathbf{w}$ with $\mathbf{x} = [1, x, x^2]^T$ and $\mathbf{w} = [w_0, w_1, w_2]^T$
- Our model is still linear in the parameters, but the actual function that is fitted is now **quadratic**:

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x + w_2 x^2$$

Polynomial Models

- We can continue adding columns of this form ...

$$\mathbf{X} = \begin{bmatrix} x_1^0 & x_1^1 & x_1^2 & \dots & x_1^K \\ x_2^0 & x_2^1 & x_2^2 & \dots & x_2^K \\ \vdots & & & & \\ x_N^0 & x_N^1 & x_N^2 & \dots & x_N^K \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- Our model function can then be written as $f(\mathbf{x}; \mathbf{w}) = \sum_{k=0}^K w_k x^k$

Coding Time!

jupyter Non_Linear_Regression Last Checkpoint: a few seconds ago (autosaved)

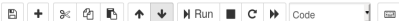


Logout

File Edit View Insert Cell Kernel Widgets Help

Trusted

Python 3



```
In [ ]: import numpy
import matplotlib.pyplot as plt
import linreg
```

```
In [ ]: # number of data points
n_points = 15

# set a seed here to initialize the random number generator
# (such that we get the same dataset each time this cell is executed)
numpy.random.seed(1)

# let's generate some "non-linear" data; note
# that the sorting step is done for visualization
# purposes only (to plot the models as connected lines)
X = numpy.random.uniform(-10,10, n_points)
t = X**2 + numpy.random.random(n_points) * 25

# reshape both arrays to make sure that we deal with
# N-dimensional Numpy arrays
t = t.reshape((len(t), 1))
X = X.reshape((len(X),1))
print("Shape of our data matrix: %s" % str(X.shape))
print("Shape of our target vector: %s" % str(t.shape))
```

```
In [ ]: # instantiate the regression model
model = linreg.LinearRegression()

# fit the model
model.fit(X,t)

# get predictions for the data points
preds = model.predict(X)
```

Arbitrary “Basis Functions”

- We can basically resort to arbitrary functions ...

$$\mathbf{x} = \begin{bmatrix} h_1(x_1) & h_2(x_1) & h_3(x_1) & \dots & h_K(x_1) \\ h_1(x_2) & h_2(x_2) & h_3(x_2) & \dots & h_K(x_2) \\ \vdots & & & & \\ h_1(x_N) & h_2(x_N) & h_3(x_N) & \dots & h_K(x_N) \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

- Our model function can then be written as $f(\mathbf{x}; \mathbf{w}) = \sum_{k=0}^K w_k h_k(x)$

Arbitrary “Basis Functions”

- We can basically resort to arbitrary functions ...

$$\mathbf{X} = \begin{bmatrix} h_1(x_1) & h_2(x_1) & h_3(x_1) & \dots & h_K(x_1) \\ h_1(x_2) & h_2(x_2) & h_3(x_2) & \dots & h_K(x_2) \\ \vdots & & & & \\ h_1(x_N) & h_2(x_N) & h_3(x_N) & \dots & h_K(x_N) \end{bmatrix} \quad \text{and} \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix}$$

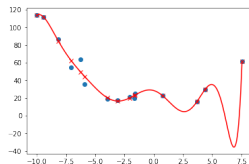
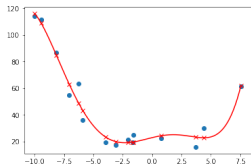
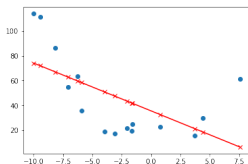
- Our model function can then be written as $f(\mathbf{x}; \mathbf{w}) = \sum_{k=0}^K w_k h_k(x)$

General Case

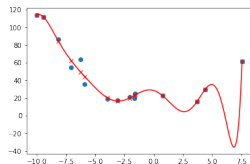
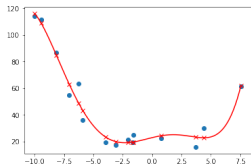
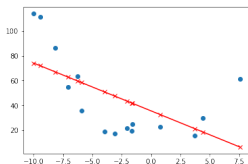
Also, given more input variables ($D > 1$), we can simply

- transform each input variable/column ...
- combine different input variables (e.g., difference between columns) ...
- combine and transform input variables ...
- ...

Which model is the best?

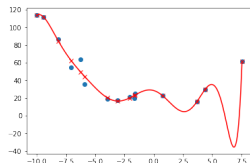
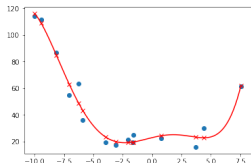
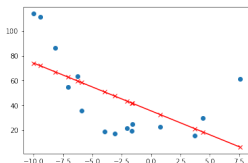


Which model is the best?



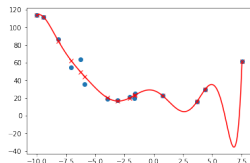
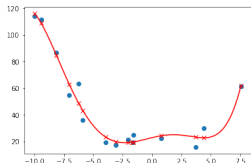
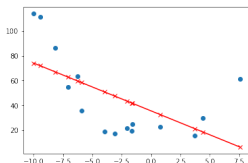
- We are given a so-called **training set** $T = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\} \subset \mathbb{R}^D \times \mathbb{R}$.

Which model is the best?



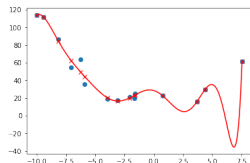
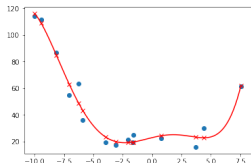
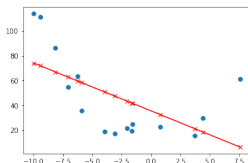
- We are given a so-called **training set** $T = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\} \subset \mathbb{R}^D \times \mathbb{R}$.
- Given the additional flexibility, we now have to **choose** a “good” model ...
 - 1 Which non-linear functions should we choose?
 - 2 How many additional columns should be generated?
 - 3 ...

Which model is the best?



- We are given a so-called **training set** $T = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\} \subset \mathbb{R}^D \times \mathbb{R}$.
- Given the additional flexibility, we now have to **choose** a “good” model ...
 - 1 Which non-linear functions should we choose?
 - 2 How many additional columns should be generated?
 - 3 ...
- We would like to choose a model that **performs well on new, unseen data!**

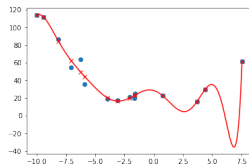
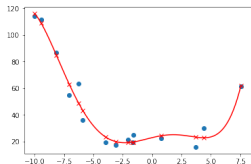
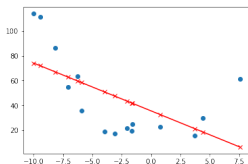
Which model is the best?



- We are given a so-called **training set** $T = \{(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N)\} \subset \mathbb{R}^D \times \mathbb{R}$.
- Given the additional flexibility, we now have to **choose** a “good” model ...
 - 1 Which non-linear functions should we choose?
 - 2 How many additional columns should be generated?
 - 3 ...
- We would like to choose a model that **performs well on new, unseen data!**

Question: How can we select such a model?

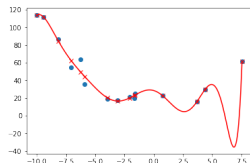
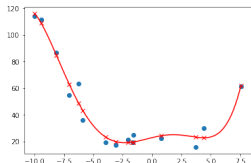
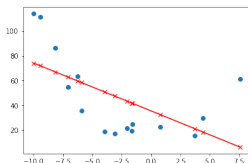
Evaluation of “Model Quality”



- The quality can be evaluated using, e.g., the **mean squared error (MSE)**:

$$\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2$$

Evaluation of “Model Quality”



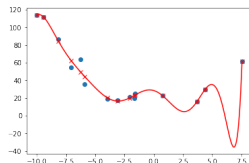
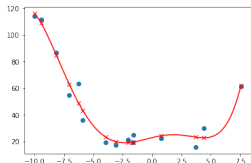
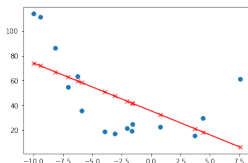
- The quality can be evaluated using, e.g., the **mean squared error (MSE)**:

$$\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2$$

- A variant is the so-called **root mean squared error (RMSE)**:

$$\sqrt{\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2}$$

Evaluation of “Model Quality”



- The quality can be evaluated using, e.g., the **mean squared error (MSE)**:

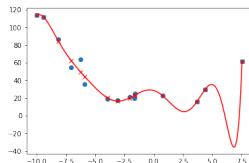
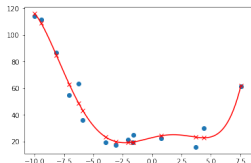
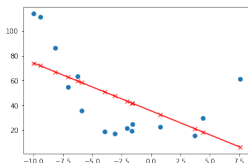
$$\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2$$

- A variant is the so-called **root mean squared error (RMSE)**:

$$\sqrt{\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2}$$

- How should we evaluate the error?

Evaluation of “Model Quality”



- The quality can be evaluated using, e.g., the **mean squared error (MSE)**:

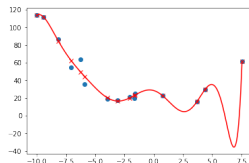
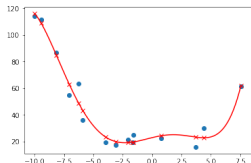
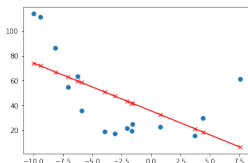
$$\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2$$

- A variant is the so-called **root mean squared error (RMSE)**:

$$\sqrt{\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2}$$

- How should we evaluate the error? On a separate dataset!

Evaluation of “Model Quality”



- The quality can be evaluated using, e.g., the **mean squared error (MSE)**:

$$\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2$$

- A variant is the so-called **root mean squared error (RMSE)**:

$$\sqrt{\frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}))^2}$$

- How should we evaluate the error? On a separate dataset! Why?

Coding Time!

jupyter Non_Linear_Regression_Training_Validation (autosaved)



Logout

File Edit View Insert Cell Kernel Widgets Help

Not Trusted

Python 3

```
In [1]: import numpy
import matplotlib.pyplot as plt
import linreg
```

```
In [2]: # set a seed here to initialize the random number generator
# (such that we get the same dataset each time this cell is executed)
numpy.random.seed(2)

# note: using large values here might lead to numerical inaccuracies
order_range = range(2,16)

# let's generate some "non-linear" data; not
# that the sorting step is done for visualization
# purposes only (to plot the models as connected lines)
X = numpy.random.uniform(-10,10,80)
t = X**2 + numpy.random.random(80) * 25

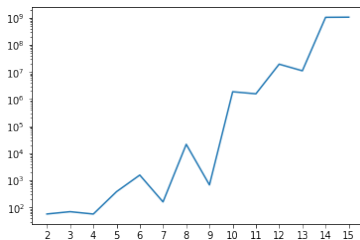
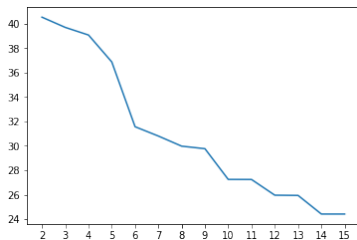
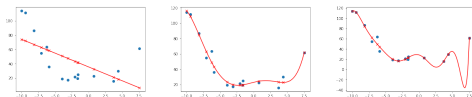
# reshape both arrays to make sure that we deal with
# N-dimensional Numpy arrays
X = X.reshape((len(X),1))
t = t.reshape((len(t), 1))
print("Shape of our data matrix: %s" % str(X.shape))
print("Shape of our target vector: %s" % str(t.shape))
```

```
Shape of our data matrix: (80, 1)
Shape of our target vector: (80, 1)
```

```
In [3]: def augment(X, max_order):
    """ Augments a given data
    matrix by adding additional
    columns.

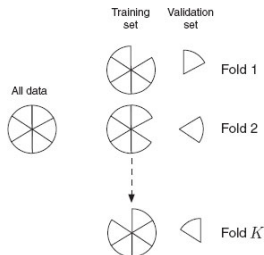
    NOTE: In case max_order is very large,
    numerical inaccuracies might occur
    """
```

Overfitting



Training error (left) and validation error (right) when fitting polynomials of increasing order (x-axis).

Training and Validation Sets



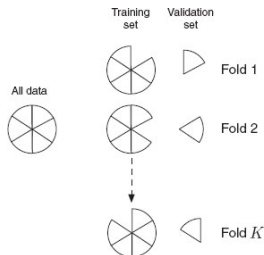
Cross-Validation

K-fold cross-validation splits the data into K (almost) equally-sized parts. We consider K “rounds”:

- 1 Use $K - 1$ parts for **training the model**. For instance, the optimal weight vector \mathbf{w} is computed for linear regression.
- 2 Use the remaining part for **validating the model** by computing the induced loss on this part.

This yields K validation errors. Typically, the average of these values is considered.

Training and Validation Sets



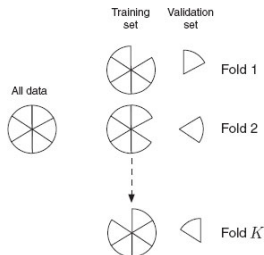
Cross-Validation

Typical values for K are $K = 2$, $K = 5$, $K = 10$, or $K = N$. The last case ($K = N$) is called **Leave-One-Out Cross Validation (LOOCV)**. The average validation error for LOOCV is given by

$$\mathcal{L}^{CV} = \frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}_{-n}))^2$$

where \mathbf{w}_{-n} is weight vector computed **without** the n -th training example.

Training and Validation Sets



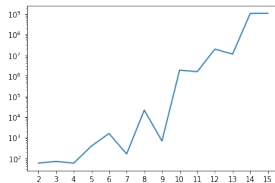
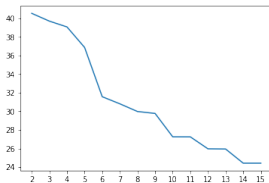
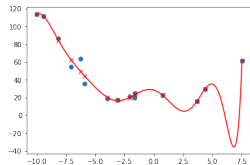
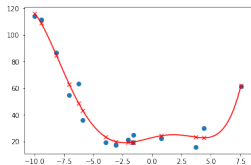
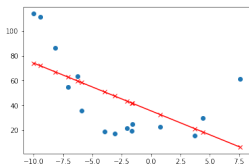
Cross-Validation

Typical values for K are $K = 2$, $K = 5$, $K = 10$, or $K = N$. The last case ($K = N$) is called **Leave-One-Out Cross Validation (LOOCV)**. The average validation error for LOOCV is given by

$$\mathcal{L}^{CV} = \frac{1}{N} \sum_{n=1}^N (t_n - f(\mathbf{x}_n; \mathbf{w}_{-n}))^2$$

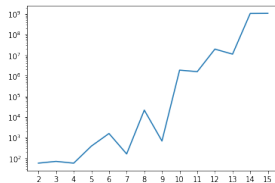
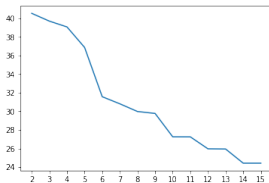
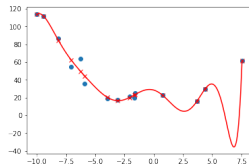
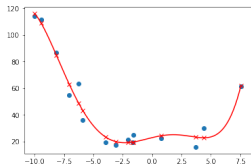
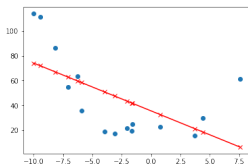
where \mathbf{w}_{-n} is weight vector computed **without** the n -th training example.

Which model is the best?



Question: How can we select such a model?

Which model is the best?



Question: How can we select such a model?

Select the model with the best validation error!

Real-World Performance

Final model performance?

Real-World Performance

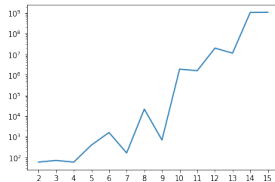
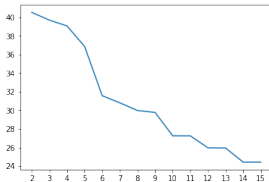
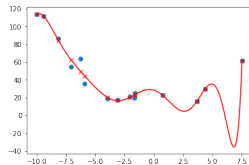
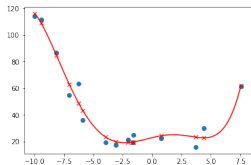
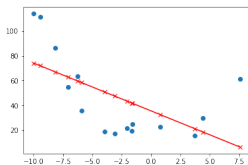
Final model performance?

*Make use of a third dataset, the so-called **test set**, for the **final** evaluation!*

Outline

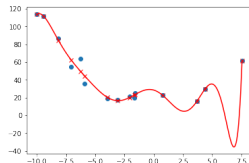
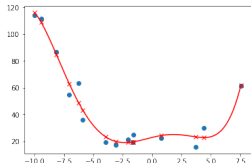
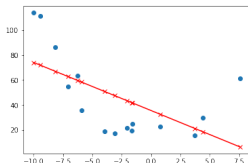
- 1 Recap: Linear Regression
- 2 Non-Linear Response, Overfitting, and Cross-Validation
- 3 Regularisation**
- 4 Summary & Outlook

Which Model is the Best?



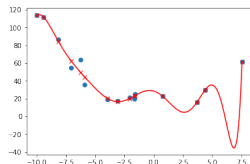
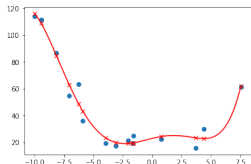
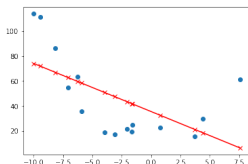
Green option: Stop increasing model complexity as soon as model becomes too complex, i.e., when it starts over-fitting (training error goes down, but validation error goes up!).

Regularisation



- The **simple** model $f(\mathbf{x}, \mathbf{w}) = \mathbf{x}^T \mathbf{w}$ with $\mathbf{w} = [0, \dots, 0]^T$ always predicts 0.

Regularisation

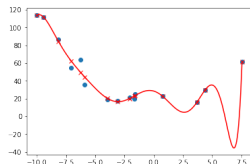
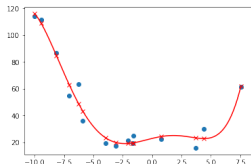
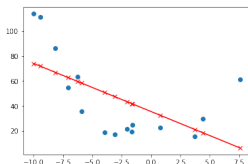


- The **simple** model $f(\mathbf{x}, \mathbf{w}) = \mathbf{x}^T \mathbf{w}$ with $\mathbf{w} = [0, \dots, 0]^T$ always predicts 0.
- Consider the following 5-th order polynomial:

$$f(x; \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + w_3 x^3 + w_4 x^4 + w_5 x^5$$

Let's start with $\mathbf{w} = \mathbf{0}$. Now, by allowing some of the w_i to be non-zero, we can make the model **more and more complex/flexible**!

Regularisation



- The **simple** model $f(\mathbf{x}, \mathbf{w}) = \mathbf{x}^T \mathbf{w}$ with $\mathbf{w} = [0, \dots, 0]^T$ always predicts 0.
- Consider the following 5-th order polynomial:

$$f(x; \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + w_3 x^3 + w_4 x^4 + w_5 x^5$$

Let's start with $\mathbf{w} = \mathbf{0}$. Now, by allowing some of the w_i to be non-zero, we can make the model **more and more complex/flexible**!

- Remember, **we don't want too complex models**! To avoid this, we can **"penalize" complex models** by adding the term $\sum_i w_i^2$ to the objective:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N}(\mathbf{X}\mathbf{w} - \mathbf{t})^T(\mathbf{X}\mathbf{w} - \mathbf{t}) + \lambda \mathbf{w}^T \mathbf{w},$$

where $\lambda > 0$ is a **model parameter** controlling the trade-off between penalising (a) not fitting the data well and (b) overly complex models.

Gradient

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} + \lambda \mathbf{w}^T \mathbf{w}$$

Gradient

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} + \lambda \mathbf{w}^T \mathbf{w}$$

Toolbox (Table 1.4 in Rogers & Girolami)

- 1 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{x} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 2 $f(\mathbf{w}) = \mathbf{x}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 3 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{w}$
- 4 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{C} \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{C} \mathbf{w}$ (if \mathbf{C} is symmetric)

We have $\nabla \mathcal{L}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w}$ and therefore:

$$\frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w} = \mathbf{0}$$

Gradient

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} + \lambda \mathbf{w}^T \mathbf{w}$$

Toolbox (Table 1.4 in Rogers & Girolami)

- 1 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{x} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 2 $f(\mathbf{w}) = \mathbf{x}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 3 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{w}$
- 4 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{C} \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{C} \mathbf{w}$ (if \mathbf{C} is symmetric)

We have $\nabla \mathcal{L}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w}$ and therefore:

$$\begin{aligned} \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w} &= \mathbf{0} \\ \Leftrightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} + N\lambda \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned}$$

Gradient

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} + \lambda \mathbf{w}^T \mathbf{w}$$

Toolbox (Table 1.4 in Rogers & Girolami)

- 1 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{x} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 2 $f(\mathbf{w}) = \mathbf{x}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 3 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{w}$
- 4 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{C} \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{C} \mathbf{w}$ (if \mathbf{C} is symmetric)

We have $\nabla \mathcal{L}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w}$ and therefore:

$$\begin{aligned} \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w} &= \mathbf{0} \\ \Leftrightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} + N\lambda \mathbf{w} &= \mathbf{X}^T \mathbf{t} \\ \Leftrightarrow (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned}$$

Gradient

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{t} + \lambda \mathbf{w}^T \mathbf{w}$$

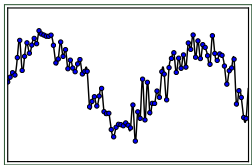
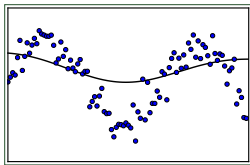
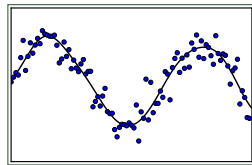
Toolbox (Table 1.4 in Rogers & Girolami)

- 1 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{x} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 2 $f(\mathbf{w}) = \mathbf{x}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = \mathbf{x}$
- 3 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{w}$
- 4 $f(\mathbf{w}) = \mathbf{w}^T \mathbf{C} \mathbf{w} \Rightarrow \nabla f(\mathbf{w}) = 2\mathbf{C} \mathbf{w}$ (if \mathbf{C} is symmetric)

We have $\nabla \mathcal{L}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w}$ and therefore:

$$\begin{aligned} \frac{2}{N} \mathbf{X}^T \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^T \mathbf{t} + 2\lambda \mathbf{w} &= \mathbf{0} \\ \Leftrightarrow \mathbf{X}^T \mathbf{X} \mathbf{w} + N\lambda \mathbf{w} &= \mathbf{X}^T \mathbf{t} \\ \Leftrightarrow (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned}$$


Regularised Linear Regression

 $\lambda = \text{small}$  $\lambda = \text{large}$  $\lambda = \text{middle}$

Minimizer for Regularised Linear Regression

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{t}$$

Coding Time!








 jupyter Non_Linear_Regression_Regularisation (unsaved changes)


[Logout](#)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted

 Python 3 

           Code 

copy selected cells

```
In [ ]: import numpy
import matplotlib.pyplot as plt
import linreg
```

```
In [ ]: # number of data points
n_points = 10

# set a seed here to initialize the random number generator
# (such that we get the same dataset each time this cell is executed)
numpy.random.seed(1)

# let's generate some "non-linear" data; note
# that the sorting step is done for visualization
# purposes only (to plot the models as connected lines)
X = numpy.random.uniform(-10,10, n_points)
t = X**2 + numpy.random.random(n_points) * 25

# reshape both arrays to make sure that we deal with
# N-dimensional Numpy arrays
t = t.reshape((len(t), 1))
X = X.reshape((len(X),1))
print("Shape of our data matrix: %s" % str(X.shape))
print("Shape of our target vector: %s" % str(t.shape))
```

```
In [ ]: # maximum degree
max_degree = 7

def augment(X, max_order):
    """ Augments a given data
    matrix by adding additional
    columns.

    NOTE: In case max_order is very large,
    numerical inaccuracies might occur
    """

    X_augmented = X
```

Multivariate Regularised Linear Regression

- **Given:** Pairs of the form $(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N) \in \mathbb{R}^D \times \mathbb{R}$ and parameter $\lambda > 0$.
- **Goal:** Find $(D + 1)$ -dimensional weight vector $\hat{\mathbf{w}} = [\hat{w}_0, \hat{w}_1, \dots, \hat{w}_D]^T$ that minimizes $\mathcal{L}(\mathbf{w}) = \frac{1}{N}(\mathbf{X}\mathbf{w} - \mathbf{t})^T(\mathbf{X}\mathbf{w} - \mathbf{t}) + \lambda \mathbf{w}^T \mathbf{w}$, i.e., which is a solution to

$$\begin{aligned} \nabla \mathcal{L}(\mathbf{w}) &= \mathbf{0} \\ \Leftrightarrow (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned} \tag{2}$$

Multivariate Regularised Linear Regression

- **Given:** Pairs of the form $(\mathbf{x}_1, t_1), \dots, (\mathbf{x}_N, t_N) \in \mathbb{R}^D \times \mathbb{R}$ and parameter $\lambda > 0$.
- **Goal:** Find $(D+1)$ -dimensional weight vector $\hat{\mathbf{w}} = [\hat{w}_0, \hat{w}_1, \dots, \hat{w}_D]^T$ that minimizes $\mathcal{L}(\mathbf{w}) = \frac{1}{N}(\mathbf{X}\mathbf{w} - \mathbf{t})^T(\mathbf{X}\mathbf{w} - \mathbf{t}) + \lambda \mathbf{w}^T \mathbf{w}$, i.e., which is a solution to

$$\begin{aligned} \nabla \mathcal{L}(\mathbf{w}) &= \mathbf{0} \\ \Leftrightarrow (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^T \mathbf{t} \end{aligned} \quad (2)$$

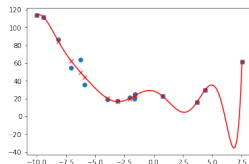
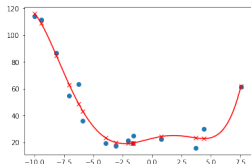
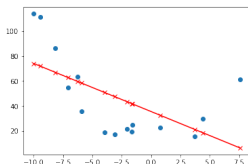
Computation in Practice

- 1 Definition of data matrix $\mathbf{X} \in \mathbb{R}^{N \times (D+1)}$
(make use of Numpy arrays and functions!)
- 2 There are different ways to compute an optimal weight vector $\hat{\mathbf{w}}$:
 - 1 Compute $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X} + N\lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{t}$ (e.g., via `numpy.linalg.inv`)
 - 2 Directly solve system of equations (2) (e.g., via `numpy.linalg.solve`)
 - 3 ...
- 3 For new point $\mathbf{x}_{new} \in \mathbb{R}^D$: Compute $t_{new} = [1, \mathbf{x}_{new}^T] \hat{\mathbf{w}}$

Outline

- 1 Recap: Linear Regression
- 2 Non-Linear Response, Overfitting, and Cross-Validation
- 3 Regularisation
- 4 Summary & Outlook**

Summary & Outlook



Today

- Recap: Multivariate linear regression
- Non-linear models by augmenting data matrix
- Cross-validation: How to select a good model ...
- Regularisation: How to avoid overfitting ...

Outlook

- Statistics (Thursday, Bulat)